

DANIEL PFELLER

SOLUTIONS ARCHITECT · GENAI / CLOUD

pfeller.daniel@gmail.com · +41 79 459 85 45 · Zürich, CH · DE (native) · EN (fluent)

*10+ years as a pre-sales technical advisor guiding large EMEA enterprises and high-growth unicorn startups through cloud and AI adoption — from use-case discovery → eval → fine-tuning → capacity planning → production. Hands-on with Claude on GCP and Gemini API across the full LLM lifecycle. Trusted by C-level executives and founders alike, currently impacting **\$800M** in cloud revenue. Committed to **safe, beneficial AI** — comfortable in both engineering deep-dives and boardroom conversations.*

SELECTED IMPACT

\$800M

cloud revenue impacted in 2025

10+ yrs

pre-sales architect — GCP, Azure, GenAI

Shipped

first Google Assistant ML model in Chrome

EXPERIENCE

Customer Engineer (Solutions Architect, Pre-Sales) — Google Cloud Platform at Google 2017 – present

- Trusted technical advisor across two customer profiles: historical book of large EMEA enterprises, currently focused on high-growth unicorn startups in scale-up phase.
- Translate business requirements into Google Cloud architectures (data platforms, GKE, Vertex AI / Gemini / Claude on Vertex, MLOps) across the full LLM lifecycle: use-case discovery → eval → fine-tuning → capacity / throughput planning → production rollout.
- Owned technical narrative end-to-end with C-suite stakeholders; impact **\$800M** in cloud revenue in 2025; recurring board-level architecture reviews and executive briefings on GenAI strategy.
- Ran use-case discovery workshops translating fuzzy business intent into a shortlist of high-value LLM use cases, scoped for measurable success and responsible deployment.
- Designed evaluation frameworks with customers: golden datasets, automated graders, side-by-side model comparison (Gemini, Claude, fine-tuned variants), regression gates for production.
- Led fine-tuning engagements on Vertex (SFT / adapter tuning) — data curation, training runs, eval, cost/quality tradeoff against base + RAG baselines.
- Sized capacity & throughput: provisioned throughput vs pay-as-you-go, latency budgets, regional placement, quota planning for enterprise launch.
- Hands-on Python prototyping with customers: notebooks, API integrations, agents, function-calling, RAG demos on customer data.
- Partnered with Google Engineering & Product Management to feed customer needs into Vertex AI / Gemini roadmap.
- Cutting-edge work with **AlphaEvolve** in partnership with Google DeepMind — bringing frontier research to enterprise customers.

Technical Program Manager — Google Assistant ML at Google 2022

- Shipped the **first Google Assistant ML model in Chrome** — owned cross-functional program across two ML engineering teams from model handoff to launch.
- Hands-on with LLMs in the pre-ChatGPT era — worked with **LaMDA** (Google's conversational model) inside Google Assistant.
- Established evaluation, monitoring, and rollback discipline for the ML release; defined OKRs and the metrics that triggered a mid-program course correction.

- Revamped coordination for an 80-engineer global program → **+20%** delivery velocity.
- Refocused scope onto highest-value user journeys → **+10%** user satisfaction.

Technology Solutions Professional (Pre-Sales) — Microsoft Azure at **Microsoft**

2014 – 2017

- Pre-sales architect through the early enterprise cloud era — drove Azure adoption for the largest accounts in the region when the category was largely uncharted.
- Designed cloud architectures and led PoCs that moved enterprises off traditional IT onto scalable cloud platforms; trusted advisor through complex enterprise buying cycles.

Technical Evangelist — Microsoft Azure at **Microsoft**

H1 2014

- Conference keynotes, deep-dive sessions, reference architectures, whitepapers, blog posts — teaching cloud-native patterns to the developer community.

Co-Founder & Software Engineer at **HaxSync (Android)**

2013 – 2014

- Built Android product reaching **100K+** downloads, top-100 paid apps, sustained 4.4★ rating — full ownership: product, engineering, support.

Teaching Fellow — Fundamentals of Computer Science at **Vienna University of Technology**

2013

- Lectures on programming, algorithms, design patterns.

Software Engineer Intern at **ETM Control GmbH**

2012

- HTML/JavaScript mobile client for a complex SCADA system.

SIDE PROJECTS — APPLIED AI IN PRODUCTION

steeze.it — LLM ski coach

ongoing

- LLM-powered ski coaching app — turns video into personalised feedback and drill plans, plus social feed for sharing runs, progress, and coaching with friends.
- End-to-end ownership: use-case framing, prompt design, eval loop, deployment, real users, full lifecycle from idea to production.

Full home & personal-life automation with LLM agents

ongoing

- Self-hosted GCP stack (Home Assistant, MQTT, InfluxDB, Grafana) wired to LLM agents (Claude, Openclaw) using Anthropic's Model Context Protocol (MCP) — daily driver for Claude API: tool use, automation flows, eval of new model releases against real use cases.

SKILLS · EDUCATION · LANGUAGES

AI / LLMs — Claude (API + Vertex), Model Context Protocol (MCP), Gemini, OSS models (Llama, Mistral, Gemma), specialized models (embeddings, rerankers, ASR, vision), RAG, agents, fine-tuning (SFT/adapters), evaluation frameworks, prompt design, capacity & throughput planning

Pre-Sales — use-case discovery, technical advisory, PoC, architecture, executive briefings, complex enterprise buying cycles

Cloud Architecture — GCP, Azure, Kubernetes, Terraform, scalable distributed systems

Programming — Python (primary), C#, Java

Leadership — cross-functional team leadership, leading without authority, teaching, mentoring

Certifications — 3 × Google Cloud Certified Professional · 4 × Microsoft Certified Solution Professional

Education — BSc Software & Information Engineering, Vienna University of Technology (2011 – 2014)

Languages — German (native) · English (fluent)